

In [2]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
```

In [3]:

```
df=pd.read_csv(r"C:\Users\pucha\Downloads\used_cars_data.csv")
df
```

Out[3]:

	S.No.	Name	Location	Year	Kilometers_Driven	Fuel_Type	Transmission	Owner_Type	Mileage
0	0	Maruti Wagon R LXI CNG	Mumbai	2010	72000	CNG	Manual	First	26.6 km/kg
1	1	Hyundai Creta 1.6 CRDi SX Option	Pune	2015	41000	Diesel	Manual	First	19.67 kmpl
2	2	Honda Jazz V	Chennai	2011	46000	Petrol	Manual	First	18.2 kmpl
3	3	Maruti Ertiga VDI	Chennai	2012	87000	Diesel	Manual	First	20.77 kmpl
4	4	Audi A4 New 2.0 TDI Multitronic	Coimbatore	2013	40670	Diesel	Automatic	Second	15.2 kmpl
...
7248	7248	Volkswagen Vento Diesel Trendline	Hyderabad	2011	89411	Diesel	Manual	First	20.54 kmpl
7249	7249	Volkswagen Polo GT TSI	Mumbai	2015	59000	Petrol	Automatic	First	17.21 kmpl
7250	7250	Nissan Micra Diesel XV	Kolkata	2012	28000	Diesel	Manual	First	23.08 kmpl
7251	7251	Volkswagen Polo GT TSI	Pune	2013	52262	Petrol	Automatic	Third	17.2 kmpl
7252	7252	Mercedes-Benz E-Class 2009-2013 E 220 CDI Avan...	Kochi	2014	72443	Diesel	Automatic	First	10.0 kmpl

7253 rows × 14 columns



In [4]:

```
pd.set_option('display.max_rows',10000000000)
pd.set_option('display.max_columns',10000000000)
pd.set_option('display.width',95)
```

In [5]:

```
print('The DataFrame has %d Rows and %d Columns'%(df.shape))
```

The DataFrame has 7253 Rows and 14 Columns

In [6]:

```
df=df[['Year','Price']]
df.columns=['year','price']
df.head(10)
```

Out[6]:

	year	price
0	2010	1.75
1	2015	12.50
2	2011	4.50
3	2012	6.00
4	2013	17.74
5	2012	2.35
6	2013	3.50
7	2016	17.50
8	2013	5.20
9	2012	1.95

In [7]:

```
df.head()
```

Out[7]:

	year	price
0	2010	1.75
1	2015	12.50
2	2011	4.50
3	2012	6.00
4	2013	17.74

In [8]:

```
df.describe()
```

Out[8]:

	year	price
count	7253.000000	6019.000000
mean	2013.365366	9.479468
std	3.254421	11.187917
min	1996.000000	0.440000
25%	2011.000000	3.500000
50%	2014.000000	5.640000
75%	2016.000000	9.950000
max	2019.000000	160.000000

In [9]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7253 entries, 0 to 7252
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   year    7253 non-null   int64  
 1   price   6019 non-null   float64
dtypes: float64(1), int64(1)
memory usage: 113.5 KB
```


In [10]:

```
sns.lmplot(x="year",y="price",data=df,order=5,ci=None)
```

C:\Users\pucha\AppData\Local\Programs\Python\Python310\lib\site-packages\seaborn\regression.py:254: RankWarning: Polyfit may be poorly conditioned
yhat = reg_func(x, y)

Out[10]:

<seaborn.axisgrid.FacetGrid at 0x2b4c4435c90>




In [11]:

```
df500=df[:,250:500]
sns.lmplot(x="year",y="price",data=df500,order=1)
```

Out[11]:

```
<seaborn.axisgrid.FacetGrid at 0x2b4c4408a00>
```



In [12]:

```
x=np.array(df['year']).reshape(-1,1)
y=np.array(df['price']).reshape(-1,1)
df.dropna(inplace=True)
```

C:\Users\pucha\AppData\Local\Temp\ipykernel_464\1387402884.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.dropna(inplace=True)
```

In [14]:

```
features_matrix=df.iloc[:,0:34]
```

In [15]:

```
target_vector=df.iloc[:,34]
```

In [16]:

```
print('The Features Matrix Has %d Rows and %d Columns'%(features_matrix.shape))  
print('The Target Matrix Has %d Rows and %d Columns'%(np.array(target_vector).reshape(-1,1).shape))
```


The Features Matrix Has 6019 Rows and 2 Columns
The Target Matrix Has 6019 Rows and 1 Columns

In [17]:

```
sns.pairplot(df)
```

Out[17]:

<seaborn.axisgrid.PairGrid at 0x2b4c4435870>




In [19]:

```
sns.displot(df['price'])
```

Out[19]:

<seaborn.axisgrid.FacetGrid at 0x2b4ca168c10>




In [20]:

```
sns.displot(df['year'])
```

Out[20]:

<seaborn.axisgrid.FacetGrid at 0x2b4c425fac0>



In [24]:

```

df500.fillna(method='ffill',inplace=True)
x=np.array(df['year']).reshape(-1,1)
y=np.array(df['price']).reshape(-1,1)
df.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='r')
plt.show()

```


C:\Users\pucha\AppData\Local\Temp\ipykernel_464\4156448208.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy (https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
df.dropna(inplace=True)
```

Regression: 0.07923369981620254



In [26]:

```

from sklearn.linear_model import LinearRegression
lm=LinearRegression()
lm.fit(x_train,y_train)

```

Out[26]:

```

LinearRegression
LinearRegression()


```

In [28]:

```
predictions=lm.predict(x_test)
plt.scatter(y_test,predictions)
```

Out[28]:

<matplotlib.collections.PathCollection at 0x2b4cc691d50>




In [29]:

```
from sklearn import metrics
print('MAE:',metrics.mean_absolute_error(y_test,predictions))
print('MSE:',metrics.mean_squared_error(y_test,predictions))
print('RMSE:',np.sqrt(metrics.mean_squared_error(y_test,predictions)))
```

MAE: 6.7706565809809565
MSE: 107.96818426543761
RMSE: 10.390773997418941

In [30]:

```
sns.displot((y_test-predictions),bins=50);
```



In [31]:

```
df.columns
```

Out[31]:

```
Index(['year', 'price'], dtype='object')
```

In [32]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2_score:",r2)
```

```
R2_score: 0.07923369981620254
```

In []: